

# HEP Software Foundation Community White Paper

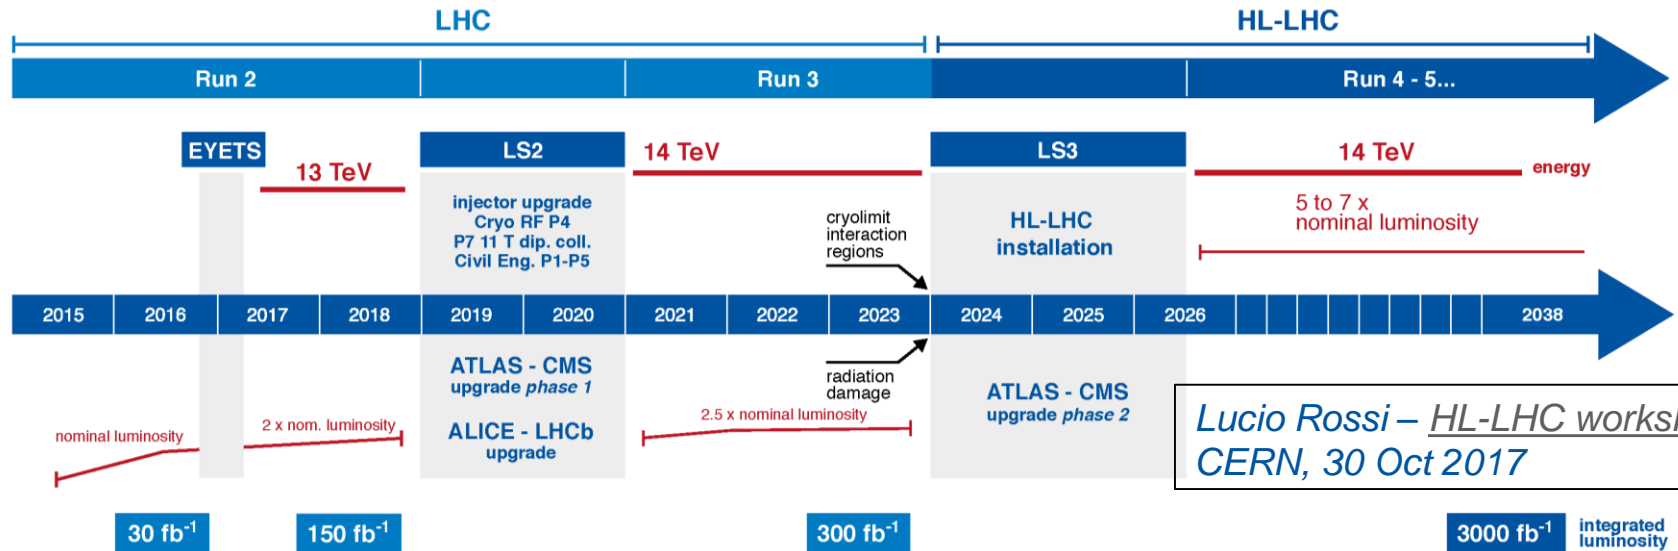
Andrea Valassi (IT-DI-LCG)

CERN IT Technical Forum – 3<sup>rd</sup> November 2017


*Disclaimer: credits to the CWP and WG authors;  
misunderstandings, comments and opinions are my own!  
(Apologies for not making it clearer what is verbatim in the CWP and what is my own.)*

# Exciting, but challenging, times ahead

## LHC / HL-LHC Plan




- Run3 (2021): major LHCb/ALICE upgrades
- Run4 (2026): major ATLAS/CMS upgrades, **high-luminosity LHC**
  - **more data** (higher luminosity) – and correspondingly more MC to generate/simulate
  - **more complex events** (higher pileup)



Pileup (2017):  
 $\mu \sim 30$  to  $60$

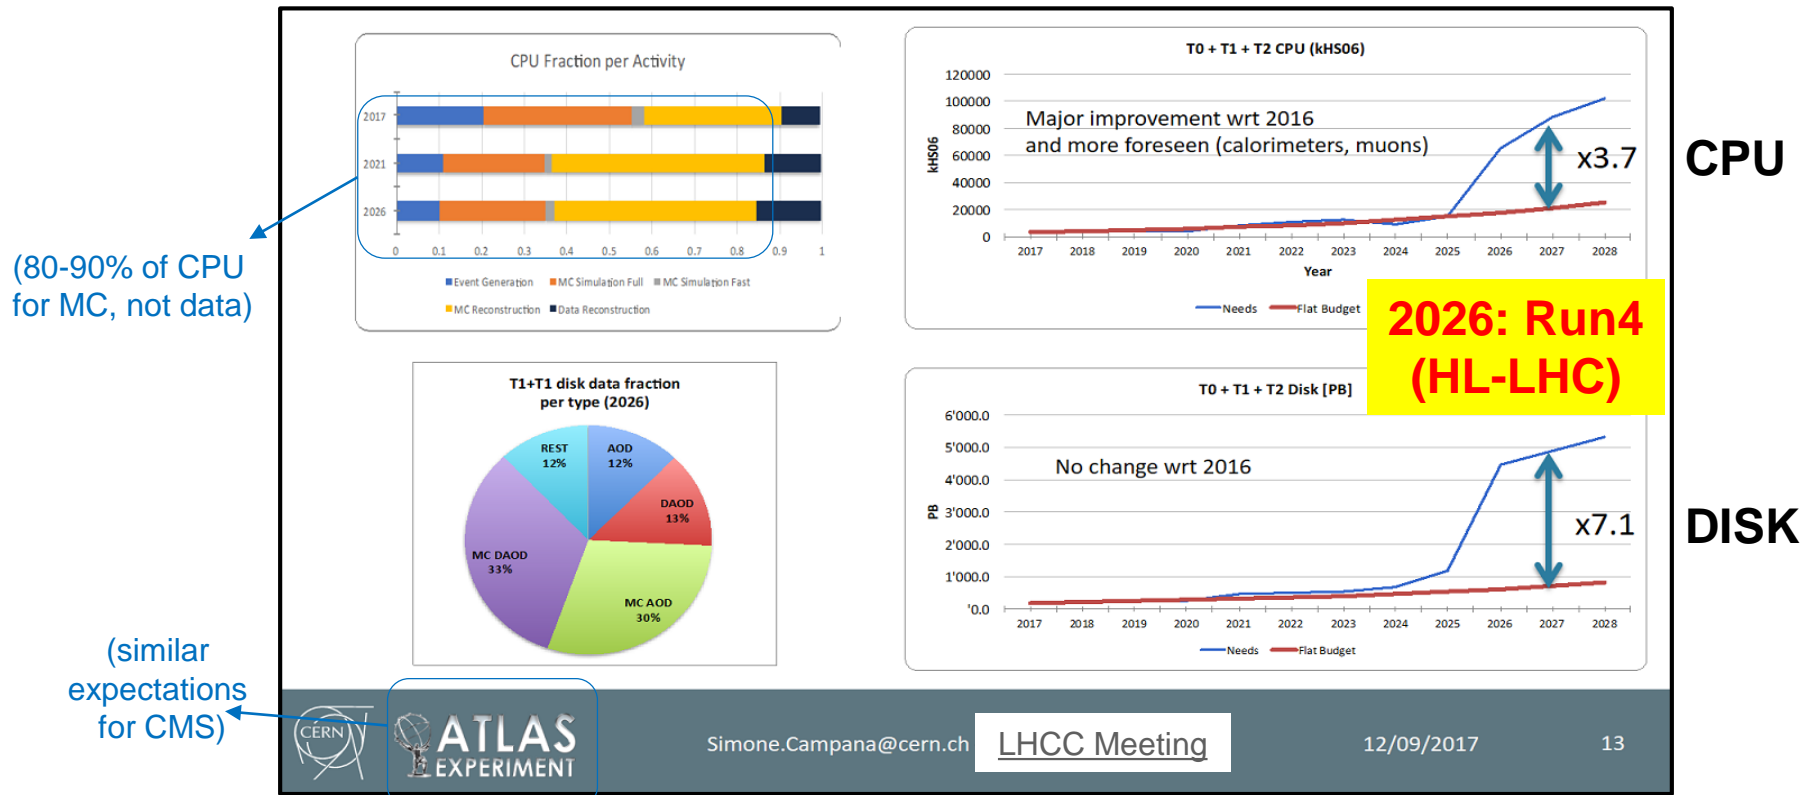
Pileup (2026):  
 $\mu \sim 130$  to  $200$

## Proof of Concept, Proof of Challenge



- Real-life event with HL-LHC-like pileup from special run in 2016 with individual high intensity bunches



# Computing resource challenges



Shortfall of resources for HL-LHC: factor ~7 for disk, factor ~4 for CPU

**“The amount of data that experiments can collect and process in the future will be limited by affordable software and computing, not by physics”**

# Technology and market trends



B. Panzer – [Technology tracking twiki](#) – May 2017

- Improvement per year: ~10-20% for CPU and ~20-30% for disk
  - already included in ATLAS and CMS projections at the LHCC
  - Moore's law and Kryder's law are slowing down
- More diverse landscape, too
  - multi- and many-core processors
  - wide vector registries
  - GPGPUs, FPGAs, ARM, HPCs...
  - memory bandwidth relatively low
  - *many programming models*
- *Technology alone will not solve the HL-LHC resource challenge*
- Need a *“software upgrade”*!

# Software (and human) challenges

- Current software of the 4 LHC experiments
  - over 12M lines of code written over the last 15 years
  - mostly written for a single architecture (x86)
  - mostly written with serial processing in mind
  - mostly experiment specific (few exceptions, e.g. ROOT or GEANT4)
    - often HEP-specific, rather than using industry standards
  - often lacking tests and quality assurance
  - often ill documented
  - often legacy and poorly maintained (e.g. with authors no longer in HEP)
  - written by many developers with varying degrees of expertise
- *This makes a software upgrade more necessary and complex!*
  - must use new programming paradigms (concurrency, vectorization...)
  - must use new architectures and new types of resources
  - must try out new ideas and new algorithms
  - (while at the same time not disrupting the running systems)

# HSF CWP inception and goals

- CWP first proposed during the LAL HSF workshop (May 2016)
  - Community roadmap for HEP S&C – basis to pursue funds and projects
- WLCG charge for producing an HSF CWP (July 2016)
  - Overall strategy and roadmap for the required “software upgrade”
- *Prioritise software research and development needed to:*
  - achieve improvements in software efficiency and performance*
  - enable new approaches in S&C to extend physics reach of the detectors*
  - ensure the long term sustainability of the software*
- Ultimate focus should be HL-LHC (Run4 and beyond)
  - Roadmap should also identify elements to be tested and used in Run3
  - Specific contact also made with FNAL muon and neutrino experiments, Belle II, Linear Collider community, national computing organisations
    - a primary goal of HSF is fostering communication across HEP and beyond



# HSF CWP in a wider context

- NSF S2I2-HEP
  - Conceptualization of a Scientific Software Innovation Institute for HEP
  - Significant contributions to the organization of HSF CWP workshops
- CERN Openlab White Paper published in September 2017
  - “Future ICT Challenges in Scientific Research”
  - Openlab also involved in the CWP process and organization
- Upcoming WLCG strategy document paper in Autumn 2017
  - Requested by LHCC, before the TDR for LHC Computing due in 2020
  - Will be largely based on HSF CWP, focusing more on short-term R&D



# Kick-off workshop – San Diego, January 2017



- HSF Workshop, SCSD/UCSD, 23-26 January 2017
  - 112 participants (86 US, 26 Europe/Asia)
- Plenaries and topical working group meetings
- Several other workshops on specific topics in the following 6 months

# CWP – WG reports and summary paper

- 13 Working Groups, producing their individual (10-50 pages) reports
  - Software Trigger and Event Reconstruction WG
  - Machine Learning WG
  - Data Access, Organisation and Management WG
  - Software Development, Deployment and Validation/Verification WG
  - Data Analysis and Interpretation WG
  - Conditions Database WG
  - Data and Software Preservation WG
  - Event Processing Frameworks WG
  - Physics Generators WG
  - Workflow and Resource Management WG
  - Visualization WG
  - Computing Models, Facilities, and Distributed Computing WG
  - Careers, Staffing and Training WG
- One overall roadmap document (“the” Community White Paper)
  - including short executive summaries (~3-5 pages each) from each WG
- See the [CWP page](#) on the HSF web site

*Reports may still be added from*

*- Security WG*

*- Workflow mgmt WG?*

# Final workshop – Annecy, June 2017



- HSF Workshop, LAPP Annecy, 26-30 June 2017
  - 90 participants (48 US, 42 Europe)
  - Reports (and further progress on writing!) from each of the 13 WGs
- The target to complete the process by August was a bit too optimistic...
  - Later decided to aim for individual WG reports by Sep and overall paper by Oct



# CWP first draft

## A Roadmap for HEP Software and Computing R&D for the 2020s

### *The HEP Software Foundation*

#### Table of Contents

<b>1 Introduction</b>	<b>2</b>
<b>2 Software and Computing Challenges</b>	<b>5</b>
<b>3 Programme of Work</b>	<b>10</b>
3.1 Conditions Databases	11
3.2 Data Analysis and Interpretation	14
3.3 Data and Software Preservation to Enable Reuse	19
3.4 Data Organisation, Management and Access	22
3.5 Data-Flow Processing Framework	27
3.6 Detector Simulation	29
3.7 Facilities and Distributed Computing	34
3.8 Machine Learning	37
3.9 Physics Generators	42
3.10 Software Development, Deployment, Validation and Verification	43
3.11 Software Trigger and Event Reconstruction	47
3.12 Visualisation	52
<b>4 Training and Careers</b>	<b>55</b>
<b>5 Conclusions</b>	<b>58</b>
<b>Appendix A - List of Workshops</b>	<b>60</b>
<b>References</b>	<b>63</b>

- Most WGs have completed their individual reports (ready for arxiv)
  - Generally 3 sections: challenges, current practices, R&D programme
- Editorial team has been set up

#### Editorial Team

- Predrag Buncic (CERN) - Alice contact
- Simone Campana (CERN) - ATLAS contact
- Peter Elmer (Princeton)
- John Harvey (CERN)
- Frank Gaede (DESY) - Linear Collider contact
- Maria Girone (CERN Openlab)
- Roger Jones (Univ. of Lancaster) - UK contact
- Michel Jouvin (LAL Orsay)
- Rob Kutschke (FNAL) - FNAL experiments contact
- Dario Menasce (INFN-Milano) - INFN contact
- Mark Neubauer (U.Illinois Urbana-Champaign)
- Stefan Roiser (CERN) - LHCb contact
- Liz Sexton-Kennedy (FNAL) - CMS contact
- Mike Sokoloff (U.Cincinnati)
- Graeme Stewart (CERN, HSF)
- Jean-Roch Vlimant (Caltech)

- First CWP draft prepared by a small subset of the Editorial Team
  - Released on 20 October
  - Send your feedback!

# (1/13) Facilities and Distributing Computing

- Major challenge is finding the best configuration for facilities and computing sites that make HL-LHC computing feasible
  - Hardware cost now dominated by disk and CPU, followed by tape and network
  - Substantial differences in regional funding models
- Upcoming (and present) changes
  - *heterogeneous resources* (HPC, volunteer and cloud computing)
  - *heterogeneous computing architectures* (different CPUs, GPGPUs, FPGAs)
  - new developments in network (e.g. SDN) and storage (e.g. object stores)
    - also large increase in network traffic from other science domains (e.g. SKA)
- R&D programme
  - Understand better the relationship between costs and performance of WLCG
    - Activity starting in the *Performance and Cost Model WG* (October GDB [presentation](#))
  - Define functionalities of a federated data center concept (“*data lake*”)
  - Consolidate data mgmt, prototype site storage cache and event streaming
  - Investigate more scalable and uniform workload scheduling
  - Prototype and evaluate a quasi-interactive analysis facility

# (2/13) Data Organization, Management and Access


- Three main challenges
  - *Higher data volumes (~factor 10) and data rates* in HL-LHC era
    - Currently ~1 Exabyte (half tape, half disk) for all experiments
    - Storage and processing costs may restrict scientific output and physics reach
  - New CPU resources (*HPC, cloud*) require more dynamic data access/mgmt
  - *Machine Learning training* needs place new requirements on data access
- Currently 3 largely independent domains: data organization (e.g. ROOT), management (transfer systems, catalogs), access (e.g. rfio, xrootd)
  - Blurring boundaries, will need to move to more global optimizations (“DOMA”)
- R&D programme
  - Study event-based granularity (*event streaming* may simplify HPC integration)
  - Evaluate Big Data techniques (Spark-like, compression, column- vs row-wise)
  - Investigate *data caching* (both for reco/analysis/simulation and for ML training)
  - Exploit varied quality of service (“tactical” storage, alternatives to tape)
  - Globally optimize data access latency using fewer, larger sites (“*data lake*”)
    - Higher availability of high-speed networks allows more extensive WAN access

# (3/13) Conditions Databases

- Conditions data: non-event data required to process event data
  - Low data volume ~TBs, but high read rate ~10kHz from ~10k jobs
  - *Caching* essential, now used by Frontier (CORAL/SQL) and CVMFS (files)
    - Ultimately provided by Squid http caches in both cases
- Challenges and opportunities
  - Event streaming (e.g. HPC) prevents job-level condition access optimizations
  - Long-term maintenance: simplify requirements (too many use cases in COOL)
  - More frequent condition updates in Run3 (decreasing instantaneous luminosity)
  - Online full reconstruction reduces the need to access conditions data offline
- R&D programme
  - New *REST-based system* (CREST) for ATLAS, in collaboration with CMS
    - LHCb have moved to a git-based conditions database
  - Identify industry-standard technologies to replace HEP-specific components
  - Study how to leverage *advanced CVMFS features* for analysis jobs



# (4/13) Detector Simulation




- Presently this means GEANT4 for most experiments
  - Advantage of a *common toolkit* is that improvements benefit everyone at once!
- *Large fraction of CPU time*: ATLAS 35% now, 25% HL-LHC
  - Including(?) the simulation of new detectors from Run4 major detector upgrades
- Main R&D goal: *speed up simulation, improve performance and efficiency*
  - Especially: efficiency on new architectures (many-core, GPGPUs, SIMD, HPC)
  - *GeantV*: ambitious R&D, multi-threaded vectorised particle transport engine
    - beta release planned for 2019, early integration of beta in the experiments for 2020
    - vectorised geometry package (VecGeom) successfully integrated in Geant4 already
  - Also explore different *fast simulation* options, including some based on ML
- Other R&D areas
  - improve physics models to achieve higher physics precision
  - investigate new options for simulating pileup (real zero-bias events, pre-mixing)
  - new pseudo-random number generator, reproducible on parallel infrastructure

# (5/13) Physics Generators


(Work in progress - section omitted from the CWP first draft)

(I am therefore describing my own – limited – understanding)



- Many generators with different features (e.g. Sherpa, MadGraph)
  - The very first step of simulation, prior to and independent of detector simulation
- *Large fraction of CPU time*: ATLAS 20% now (20B events), 10% HL-LHC
- (Some of the) main challenges for the future in my understanding
  - Need *more precise MC generators* (NNLO i.e. next-to-next-to-leading-order)
    - A very active area of theoretical physics research at the moment
  - These will be *even more computationally intensive* than current ones (NLO)
    - And will require more specialized computing such as massively parallel HPCs
  - Unclear to me if the needed resources have been (or can presently be) predicted
    - This may well be another high-priority area of work for software/computing
- See talks by [F. Maltoni](#) and [G. Zanderighi](#) at the recent HL-LHC Workshop

# (6/13) Software Trigger and Event Reconstruction



- *Largest fraction of CPU time* is event reconstruction
  - ATLAS 40% now, 60% HL-LHC (for MC ~3x as for real data)
    - Dominated by charged particle tracking (less by calorimeter, particle flow, particle ID)
  - Including the effect of higher pileup on the complexity of individual events
  - Including(?) event reconstruction in new detectors from Run4 major upgrades
  - Need benchmarks to measure *physics-per-EUR* and/or per-watt
- Major trigger upgrades upcoming
  - Run4 ATLAS/CMS: factor 10 increase, stream 1 MHz from hardware trigger
  - Run3 LHCb/ALICE: full bx rate (LHCb 40 MHz) to ~real-time *software triggers*
- R&D programme
  - Improve use of *vector units (SIMD)*
  - More efficient use of *many-core parallelism* by using multi-threaded frameworks
  - Investigate offload to dedicated heterogeneous resources (*GPGPUs, FPGAs*)
  - Improve on QA, QC, *testing and physics validation* (also using CI)
  - Evaluate tools and techniques to facilitate *real-time analysis*
  - Physics performance (also at low  $p_T$ ) with high pileup, add timing information
  - Investigate radically new algorithms, also based on ML techniques

# (7/13) Data Processing Frameworks

- *Many common concepts, but several different implementations*
  - Gaudi developed by LHCb, but also used by ATLAS (Athena)
  - CMSSW developed by CMS, forked by Intensity Frontier experiments (art)
  - Basf2 used by Belle II
  - FairROOT and AliROOT closely related, ALICE and FAIR now developing O2
- Key idea to promote in this area is *commonality*
  - Vital to maintain cross-project communication (e.g. Concurrency Forum)
    - Many new ideas to try out, including concurrency and functional programming
  - If possible, try to encapsulate new developments as *common libraries*
    - So that they could be used from different experiment frameworks
    - IMO, “framework consolidation between experiments” beyond this is highly unlikely
  - IMO, frameworks also important to provide *commonality within each experiment*
    - Let the experts do the complex stuff once for all users in a collaboration (more later)

# (8/13) Training and Careers (CWP and – a lot – my opinions)

- *Human challenge*, in addition to resource and software challenges:
  - many developers with varying expertise in a collaboration
    - specific to HEP? – users (physicists) are developers, overlap with computing experts
    - solutions evolve continuously based on what has been observed
  - high turn around, people often leaving the field
    - “we” are left with their code, “they” are left with what they learnt in HEP (reusable?)
- Solutions include
  - *training*, share experts’ knowledge with non-experts, spread best practices
    - nice examples: LHCb starter kits (inspired by Software Carpentry) and hackathons
    - computing schools (CSC, Bertinoro, GridKa) are useful but not enough
    - experiment with different tools, including notebooks, WikiToLearn, MOOCs
  - encapsulating new paradigms in *framework libraries and interfaces*
    - let experts work out the technicalities and offer simple solutions to non-experts
  - developing *common libraries and projects* with other experiments
    - a common cross-experiment Q&A (like StackExchange) would have been nice
  - establishing links with *experts from other fields*
  - *improving recognition* of software (and training) work for academic careers
    - passes through proper journal publication, too

# (9/13) Software Development practices

- Challenges when  $O(100s)$  people in a collaboration develop code
  - Code quality, modularity, reusability, maintainability
  - Code validation (from unit tests to large scale physics performance tests)
  - Distributed development between geographically separated peers
  - Deployment and interaction with operation teams
  - Software licensing and distribution (and their impact on interoperability)
  - Inhomogeneous community with varying skills and funding models
    - Additional activity splits: production vs. user analysis, online vs. offline
- Some best practices that are used or are being adopted
  - Well defined dev environments, *support multiple platforms*
  - Moving to non-HEP-specific tools (git, CMake)
  - Promoting *code testing* (in advance of MRs) and *continuous integration*
- Ongoing and proposed R&D goals (in addition to those in “Training” and “Frameworks”)
  - Continue with HSF working group on packaging
  - Strengthen links to communities outside HEP
  - Develop *tools to measure software performance* and collect profiling info
  - Common *toolkits for statistical regression testing* of reco/simu software

# (10/13) Data and Software Preservation

- Must *preserve the data and the knowledge* needed to derive results from it
  - Low-level bit preservation of the data is largely covered by DPHEP group
  - *Knowledge preservation (software, workflows, documentation) far from easy*
- Several benefits and use cases
  - Analyse old data, to resolve conflicts between results or test new theories
    - e.g. ATLAS exploring ability to reinterpret searches using services like RECAST
  - Ease transfer of knowledge within a collaboration – ensure *reproducibility*
  - Expose data for outreach purposes or to engage external experts (e.g. in ML)
    - All LHC experiments have released some of their data in the *Open Data Portal*
- R&D goals
  - Prototype analysis ecosystems including *analysis portals with a working UI*
  - Evaluate limits of *container technology* in the preservation arena
  - Demonstrate ability to execute production workflows including many containers
  - Collect analysis use cases, tracking evolution towards Big Data environments



# (11/13) Machine Learning

- HEP has used “ML” (e.g. neural networks to classify signal/bkg) for >20yrs
  - But now ML is ubiquitous in Big Data and many mature open source tools exist
    - BDT/NN are now in most HEP papers too, e.g. for pID or b/c/light jet classification
    - HEP-specific (TMVA in ROOT) and non-HEP (Scikit-learn, Keras...) packages
  - Challenge: *new, faster ML algos* for track/vertex reco, trigger, simulation?
    - While maintaining “sufficient fidelity” to the raw data
  - Other challenges from ML computing infrastructure (training, interactive UI...)
- R&D programme (coordinated by common *Inter-Experiment ML Forum*)
  - *Evaluation of data formats* (ROOT vs others) and impact on ML performance
  - Explore *ML as a Service (MLaS)*, starting from what is now provided by *SWAN*
  - Try out Deep Learning computer-vision techniques, e.g. for pID in LArTPCs
  - Evaluate ML for network anomaly detection to optimize data transfers
  - And many other ideas for simulation, trigger, tracking, anomaly detection...

# (12/13) Data Analysis and Interpretation

- Final stage of analysis – e.g. final data reduction and interactive plots
  - not part of “production” computing (e.g. not in ATLAS LHCC plots) – laptops, too
  - HEP currently gravitates around *ROOT* (I/O, ntuples, histograms, plots...)
    - challenges: *maintenance and sustainability* of a HEP-specific product
  - rise of many alternatives outside HEP (Python ecosystem, Spark...)
    - challenges: *interchangeability and interoperability* (“bridges” and “ferries”)
- [R&D programme](#) (largely discussed also at the [Amsterdam workshop](#))
  - *Make Python a first-class citizen* in HEP (specifically in ROOT)
  - Enable dynamically *plugging in open-source tools* (interchangeability)
  - Interoperability: *bridges* (use external tools) and *ferries* (use external data)
  - Prototype *functional and declarative* programming models
  - *Analysis facilities* (e.g. Spark-like, query-based) beyond current “primitive” ones
  - Prototype an Interpretation Gateway, extending present recasting tools

# (13/13) Visualisation

- First main use case: *event displays* and *geometry visualisation*
  - 3d rendering via OpenGL (standalone) or WebGL (browsers)
  - Many experiment-specific tools
- Second main use case: statistical visualisation (e.g. histograms)
  - *ROOT* widely used in HEP
  - Increasing use of *Python (Matplotlib)* and JavaScript – e.g. *Jupyter notebooks*
- Common feature is *interactivity*
- Main challenges and R&D goals
  - Code *maintainability* and sustainability
  - Leverage *common packages* to decrease experiment-dependency
  - Improve access to geometry and event data through a streaming interface

# CWP - Conclusions

- ATLAS and CMS needs at HL-LHC exceed those achievable by incremental changes to code and computing facilities within the foreseen budget
- HEP S&C needs a *step change* and an *investment in people* to solve these issues
- Main goals should be
  - improvements in software (and computing) efficiency, scalability and performance
  - enable new approaches that can radically extend the physics reach of the detectors
  - ensure the long term sustainability of the software
- We must work together as a community and link to experts outside of HEP
  - “Together” includes all four LHC experiments, Belle II, DUNE, ILC...
  - Challenges will start sooner for LHCb, ALICE (and others) than for ATLAS and CMS

**“The amount of data that experiments can collect and process in the future will be limited by affordable software and computing, not by physics”**

# CWP and HSF outlook

- Second draft of the CWP expected fairly soon
  - Work still ongoing in some WGs and comments being received
  - Close 1<sup>st</sup> draft for comments on Nov 10? (then release 2<sup>nd</sup> draft)
- Other HSF activities are going on in parallel
  - Licensing, packaging, training...
  - But the emphasis in recent months has clearly been on the CWP
- Joint HSF-WLCG Workshop in Napoli (26-29 March 2018)
  - Bring together both the computing and software communities
  - The HL-LHC challenges must be addressed on both sides together!
- HSF session at CHEP2018 (probably on Sat-Sun 7-8 July )

# Backup slides

# HL-LHC expected shortfall of resources

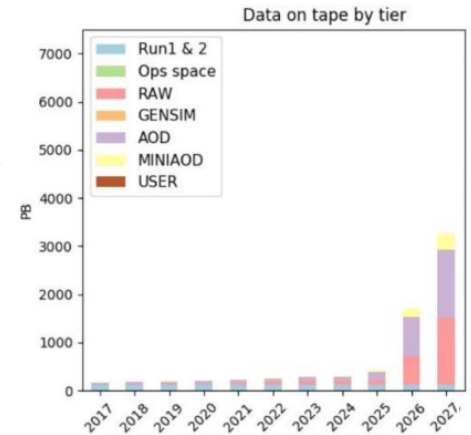
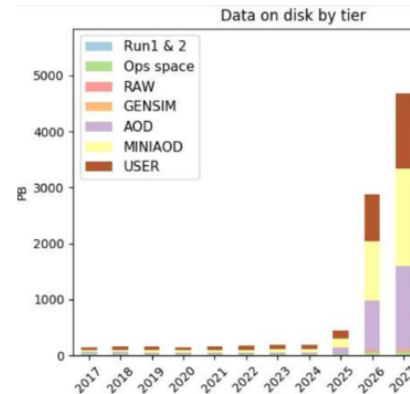
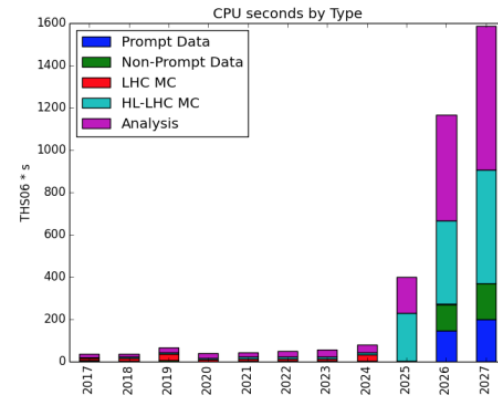
## CMS

### HL-LHC new working numbers

- ▶ CMS does not have newer officially blessed numbers for HL-LHC
- ▶ Still, work has been ongoing also due to the DOE request to US-CMS for long time planning
- ▶ Main changes wrt to older models (see for example ECFA presentation by S.Campana) are
  - ▶ Expectation of 10%/y code performance improvement
  - ▶ Rely largely on MiniAOD(SIM) for operations; AOD(SIM) an archival thing

#### Take home messages for 2027:

- 50 MHS06 CPU
- 5EB disk
- 3EB tape
- Wrt to 2017, assuming a +20%/y by Moore and friends, the excesses are ~6x for CPU, ~4x for storage



Liz Sexton-Kennedy – LHCC Meeting – 12 Sep 2017



# Run3 and Run4 upgrades



## Experimental upgrades

### Preparation in LS2 for Run 3

#### LHCb

- Significant detector upgrade
  - Rebuild of all tracking detectors
- Triggerless readout @ 40MHz
- Trigger fully implemented in software

#### ALICE:

- Significant detector upgrade (inner tracking, TPC)
- Triggerless readout for Pb-Pb @ 50kHz
- New integrated DAQ-online system with smart data compression

### Preparation in LS3 for Run 4

#### CMS

- New silicon tracker with trigger capabilities (at 1A rate)
- Improved L1 latency and calorimeter end-caps

#### ATLAS:

- New inner tracker and readout improvements in calorimeters
- Track trigger operating at level-0 rates

ACAT 2016

22. January 2016

Johannes Albrecht

4/40



# LO, NLO and NNLO

► Perturbative QCD: expansion in powers of  $\alpha_s \ll 1$

► Expansion of observable


$$f = f_0 + \alpha_s f_1 + \alpha_s^2 f_2 + \alpha_s^3 f_3 + \dots$$

► often compute only the first one (leading order, LO) or two (next-to-leading order, NLO) terms

► Technique for calculations : using Feynman diagrams

Aude Gehrmann-De Ridder

Perturbative QCD - Lecture



CERN Academic Training 2013

Precision of NLO (about 10-50%) is in many cases not enough to match current experimental precision, let alone the projected one at  $3\text{ab}^{-1}$

Giulia Zanderighi  
CERN Theory Department, University of Oxford and ERC  
High-Luminosity workshop, CERN, November 2017

HL-LHC Workshop Nov 2017